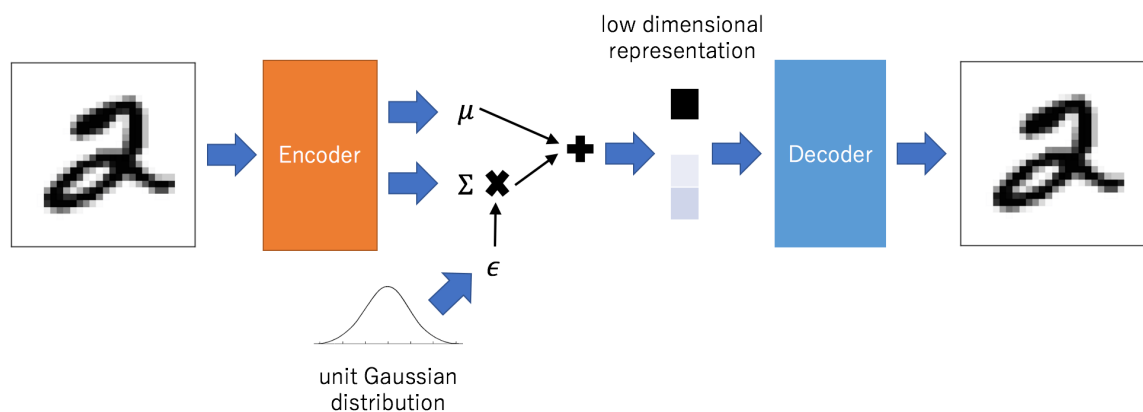


# Problem 1 - Variational Auto-Encoder (VAE)



Variational Auto-Encoders (VAEs) are a widely used class of generative models. They are simple to implement and, in contrast to other generative model classes like Generative Adversarial Networks, they optimize an explicit maximum likelihood objective to train the model. Finally, their architecture makes them well-suited for unsupervised representation learning, i.e. learning low-dimensional representations of high-dimensional inputs, like images, with only self-supervised objectives (data reconstruction in the case of VAEs).



(image source: <https://mlexplained.com/2017/12/28/an-intuitive-explanation-of-variational-autoencoders-vaes-part-1>)

**By working on this problem you will learn and practice the following steps:**

1. Set up a data loading pipeline in PyTorch.
2. Implement, train and visualize an auto-encoder architecture.
3. Extend your implementation to a variational auto-encoder.
4. Learn how to tune the critical beta parameter of your VAE.
5. Inspect the learned representation of your VAE.

**Note:** For faster training of the models in this assignment you can use Colab with enabled GPU support. In Colab, navigate to "Runtime" --> "Change Runtime Type" and set the "Hardware Accelerator" to "GPU".

## 1. MNIST Dataset

We will perform all experiments for this problem using the [MNIST dataset](#), a standard dataset of handwritten digits. The main benefits of this dataset are that it is small and relatively easy to model. It therefore allows for quick experimentation and serves as initial test bed in many papers.

Another benefit is that it is so widely used that PyTorch even provides functionality to

automatically download it.

Let's start by downloading the data and visualizing some samples.

```
In [ ]: import matplotlib.pyplot as plt
import matplotlib inline

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-i
%load_ext autoreload
%autoreload 2
```

```
In [ ]: import torch
import torchvision

# this will automatically download the MNIST training set
mnist_train = torchvision.datasets.MNIST(root='./data',
                                         train=True,
                                         download=True,
                                         transform=torchvision.transforms.T
print("\n Download complete! Downloaded {} training examples!".format(len(m

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to
./data/MNIST/raw/train-images-idx3-ubyte.gz

Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to
./data/MNIST/raw/train-labels-idx1-ubyte.gz

Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to
./data/MNIST/raw/t10k-images-idx3-ubyte.gz

Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to
./data/MNIST/raw/t10k-labels-idx1-ubyte.gz


Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

Download complete! Downloaded 60000 training examples!
```

```
In [ ]: import matplotlib.pyplot as plt
import numpy as np

# Let's display some of the training samples.
sample_images = []
mnist_it = iter(mnist_train) # create simple iterator, later we will use p
for _ in range(5):
    sample = next(mnist_it) # samples a tuple (image, label)
    sample_images.append(sample[0][0].data.cpu().numpy())

fig = plt.figure(figsize = (10, 50))
ax1 = plt.subplot(111)
ax1.imshow(np.concatenate(sample_images, axis=1), cmap='gray')
plt.show()
```



## 2. Auto-Encoder

Before implementing the full VAE, we will first implement an **auto-encoder architecture**.

Auto-encoders feature the same encoder-decoder architecture as VAEs and therefore also learn a low-dimensional representation of the input data without supervision. In contrast to VAEs they are **fully deterministic** models and do not employ variational inference for optimization.

The **architecture** is very simple: we will encode the input image into a low-dimensional representation using a convolutional network with strided convolutions that reduce the image resolution in every layer. This results in a low-dimensional representation of the input image. This representation will get decoded back into the dimensionality of the input image using a convolutional decoder network that mirrors the architecture of the encoder. It employs transposed convolutions to increase the resolution of its input in every layer. The whole model is trained by **minimizing a reconstruction loss** between the input and the decoded image.

Intuitively, the **auto-encoder needs to compress the information contained in the input image** into a much lower dimensional representation (e.g.  $28 \times 28 = 784$ px vs. 64 embedding dimensions for our MNIST model). This is possible since the information captured in the pixels is *highly redundant*. E.g. encoding an MNIST image requires  $<4$  bits to encode which of the 10 possible digits is displayed and a few additional bits to capture information about shape and orientation. This is much less than the  $255^{28 \cdot 28}$  bits of information that could be theoretically captured in the input image.

Learning such a **compressed representation can make downstream task learning easier**. For example, learning to add two numbers based on the inferred digits is much easier than performing the task based on two piles of pixel values that depict the digits.

In the following, we will first define the architecture of encoder and decoder and then train the auto-encoder model.

### Defining the Auto-Encoder Architecture [6pt]

```
In [ ]: import torch.nn as nn
```

```
# Let's define encoder and decoder networks
#####
# Encoder Architecture:
# - Conv2d, hidden units: 32, output resolution: 14x14, kernel: 4 #
# - LeakyReLU
# - Conv2d, hidden units: 64, output resolution: 7x7, kernel: 4 #
# - BatchNorm2d
# - LeakyReLU
# - Conv2d, hidden units: 128, output resolution: 3x3, kernel: 3 #
# - BatchNorm2d
# - LeakyReLU
# - Conv2d, hidden units: 256, output resolution: 1x1, kernel: 3 #
# - BatchNorm2d
# - LeakyReLU
# - Flatten
# - Linear, output units: nz (= representation dimensionality) #
#####

class Encoder(nn.Module):
    def __init__(self, nz):
        super().__init__()
        ##### TODO #####
        # Create the network architecture using a nn.Sequential module wrapper.
        # All convolutional layers should also learn a bias.
        # HINT: use the given information to compute stride and padding
        #         for each convolutional layer. Verify the shapes of intermediate
        #         by running partial networks (with the next cell) and visualizing
        #         output shapes.
        #####
        self.net = nn.Sequential(
            # add your network layers here
            # ...
            nn.Conv2d(in_channels=1, out_channels=32, kernel_size=4, stride=2,
            nn.LeakyReLU(negative_slope=1e-2),
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=4, stride=2,
            nn.BatchNorm2d(64),
            nn.LeakyReLU(negative_slope=1e-2),
            nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=2,
            nn.BatchNorm2d(128),
            nn.LeakyReLU(negative_slope=1e-2),
            nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3, stride=2,
            nn.BatchNorm2d(256),
            nn.LeakyReLU(negative_slope=1e-2),
            nn.Flatten(),
            nn.Linear(in_features=256, out_features=nz, bias=True)
        )
        ##### END TODO #####

    def forward(self, x):
        return self.net(x)

#####
# Decoder Architecture (mirrors encoder architecture):
# - Linear, output units: 256
# - Reshape, output shape: (256, 1, 1)
# - BatchNorm2d
# - LeakyReLU
# - ConvT2d, hidden units: 128, output resolution: 3x3, kernel: 3 #
# - BatchNorm2d
# - LeakyReLU
```

```

# - LeakyReLU
# - ConvT2d, hidden units: 64, output resolution: 7x7, kernel: 3
# - ...
# - ...
# - ConvT2d, output units: 1, output resolution: 28x28, kernel: 4
# - Sigmoid (to limit output in range [0...1])
#####

class Decoder(nn.Module):
    def __init__(self, nz):
        super().__init__()
        ##### TODO #####
        # Create the network architecture using a nn.Sequential module wrapper.
        # Again, all (transposed) convolutional layers should also learn a bias
        # We need to separate the initial linear layer into a separate variable
        # nn.Sequential does not support reshaping. Instead the "Reshape" is pe
        # in the forward() function below and does not need to be added to self
        # HINT: use the class nn.ConvTranspose2d for the transposed convolution
        # Verify the shapes of intermediate layers by running partial net
        # (using the next cell) and visualizing the output shapes.
        #####
        self.map = nn.Linear(in_features=nz, out_features=256, bias=True) # f
        self.net = nn.Sequential(
            # add your network layers here
            # ...
            nn.BatchNorm2d(256),
            nn.LeakyReLU(),
            nn.ConvTranspose2d(in_channels=256, out_channels=128, kernel_size=3,
            nn.BatchNorm2d(128),
            nn.LeakyReLU(),
            nn.ConvTranspose2d(in_channels=128, out_channels=64, kernel_size=3,
            nn.BatchNorm2d(64),
            nn.LeakyReLU(),
            nn.ConvTranspose2d(in_channels=64, out_channels=32, kernel_size=4,
            nn.LeakyReLU(),
            nn.ConvTranspose2d(in_channels=32, out_channels=1, kernel_size=4, s
            nn.Sigmoid()
        )
        ##### END TODO #####

    def forward(self, x):
        return self.net(self.map(x).reshape(-1, 256, 1, 1))

```

## Testing the Auto-Encoder Forward Pass [1pt]

```
In [ ]: # To test your encoder/decoder, let's encode/decode some sample images
# first, make a PyTorch DataLoader object to sample data batches
batch_size = 64
nworkers = 4 # number of workers used for efficient data loading

##### TODO #####
# Create a PyTorch DataLoader object for efficiently generating training batches
# Make sure that the data loader automatically shuffles the training dataset
# HINT: The DataLoader wraps the MNIST dataset class we created earlier.
# Use the given batch_size and number of data loading workers when creating
# the DataLoader.
#####
mnist_data_loader = torch.utils.data.DataLoader(mnist_train, batch_size=batch_size,
##### END TODO #####

# now we can run a forward pass for encoder and decoder and check the product
nz = 64 # dimensionality of the learned embedding
encoder = Encoder(nz)
decoder = Decoder(nz)
for sample_img, sample_label in mnist_data_loader:
    enc = encoder(sample_img)
    print("Shape of encoding vector (should be [batch_size, nz]): {}".format(enc.size()))
    dec = decoder(enc)
    print("Shape of decoded image (should be [batch_size, 1, 28, 28]): {}".format(dec.size()))
    break
```

```
/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:481:
UserWarning: This DataLoader will create 4 worker processes in total. Our suggested
max number of worker in current system is 2, which is smaller than what this
DataLoader is going to create. Please be aware that excessive worker creation
might get DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
  cpuset_checked))
Shape of encoding vector (should be [batch_size, nz]): torch.Size([64, 64])
Shape of decoded image (should be [batch_size, 1, 28, 28]): torch.Size([64,
1, 28, 28])
```

Now that we defined encoder and decoder network our architecture is nearly complete.

However, before we start training, we can wrap encoder and decoder into an auto-encoder class for easier handling.

```
In [ ]: class AutoEncoder(nn.Module):
    def __init__(self, nz):
        super().__init__()
        self.encoder = Encoder(nz)
        self.decoder = Decoder(nz)

    def forward(self, x):
        return self.decoder(self.encoder(x))

    def reconstruct(self, x):
        """Only used later for visualization."""
        return self.forward(x)
```

## Setting up the Auto-Encoder Training Loop [6pt]

After implementing the network architecture, we can now set up the training loop and run training.

```

In [ ]: epochs = 10
learning_rate = 1e-3

# build AE model
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu') #
ae_model = AutoEncoder(nz).to(device) # transfer model to GPU if available
ae_model = ae_model.train() # set model in train mode (eg batchnorm param

# build optimizer and loss function
##### TODO #####
# Create the optimizer and loss classes. For the loss you can use a loss la
# from the torch.nn package.
# HINT: We will use the Adam optimizer (learning rate given above, otherwis
# default parameters) and MSE loss for the criterion / loss.
# NOTE: We could also use alternative loss functions like cross entropy, de
# on the assumptions we are making about the output distribution. Her
# will use MSE loss as it is the most common choice, assuming a Gauss
# output distribution.
#####
opt = torch.optim.Adam(params=ae_model.parameters(), lr=learning_rate)
criterion = nn.MSELoss() # create loss layer instance
##### END TODO #####

train_it = 0
for ep in range(epochs):
    print("Run Epoch {}".format(ep))
    ##### TODO #####
    # Implement the main training loop for the auto-encoder model.
    # HINT: Your training loop should sample batches from the data loader, ru
    # forward pass of the AE, compute the loss, perform the backward pa
    # perform one gradient step with the optimizer.
    # HINT: Don't forget to erase old gradients before performing the backwar
    #####
    for sample_img, sample_label in mnist_data_loader:
        # add training loop commands here
        # ...
        # Move data to GPU
        sample_img = sample_img.to(device)

        opt.zero_grad()
        target = ae_model.forward(sample_img)
        rec_loss = criterion(sample_img, target)
        rec_loss.backward()
        opt.step()
        ##### END TODO #####

        if train_it % 100 == 0:
            print("It {}: Reconstruction Loss: {}".format(train_it, rec_loss))
            train_it += 1

    print("Done!")

```

Run Epoch 0

/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:481:  
 UserWarning: This DataLoader will create 4 worker processes in total. Our s  
 uggested max number of worker in current system is 2, which is smaller than  
 what this DataLoader is going to create. Please be aware that excessive wor  
 ker creation might get DataLoader running slow or even freeze, lower the wo  
 rker number to avoid potential slowness/freeze if necessary.

cpuset\_checked))

It 0: Reconstruction Loss: 0.18664182722568512

It 100: Reconstruction Loss: 0.0263246800750494

It 200: Reconstruction Loss: 0.016419680789113045  
It 300: Reconstruction Loss: 0.013837716542184353  
It 400: Reconstruction Loss: 0.011019397526979446  
It 500: Reconstruction Loss: 0.011521183885633945  
It 600: Reconstruction Loss: 0.009608311578631401  
It 700: Reconstruction Loss: 0.009960976429283619  
It 800: Reconstruction Loss: 0.009381072595715523  
It 900: Reconstruction Loss: 0.008728671818971634  
Run Epoch 1  
It 1000: Reconstruction Loss: 0.008005676791071892  
It 1100: Reconstruction Loss: 0.007739451713860035  
It 1200: Reconstruction Loss: 0.007920853793621063  
It 1300: Reconstruction Loss: 0.008873661048710346  
It 1400: Reconstruction Loss: 0.006991246249526739  
It 1500: Reconstruction Loss: 0.007462984416633844  
It 1600: Reconstruction Loss: 0.0065509844571352005  
It 1700: Reconstruction Loss: 0.007878832519054413  
It 1800: Reconstruction Loss: 0.007533562369644642  
Run Epoch 2  
It 1900: Reconstruction Loss: 0.006996598560363054  
It 2000: Reconstruction Loss: 0.006903706584125757  
It 2100: Reconstruction Loss: 0.005974492523819208  
It 2200: Reconstruction Loss: 0.006070496514439583  
It 2300: Reconstruction Loss: 0.0068871863186359406  
It 2400: Reconstruction Loss: 0.007140104193240404  
It 2500: Reconstruction Loss: 0.006445891223847866  
It 2600: Reconstruction Loss: 0.005407628137618303  
It 2700: Reconstruction Loss: 0.005920184776186943  
It 2800: Reconstruction Loss: 0.0068471599370241165  
Run Epoch 3  
It 2900: Reconstruction Loss: 0.005907618906348944  
It 3000: Reconstruction Loss: 0.007192608900368214  
It 3100: Reconstruction Loss: 0.00592842185869813  
It 3200: Reconstruction Loss: 0.006346351932734251  
It 3300: Reconstruction Loss: 0.00574900396168232  
It 3400: Reconstruction Loss: 0.005313708446919918  
It 3500: Reconstruction Loss: 0.005397276021540165  
It 3600: Reconstruction Loss: 0.005362116266041994  
It 3700: Reconstruction Loss: 0.005842178128659725  
Run Epoch 4  
It 3800: Reconstruction Loss: 0.005256583448499441  
It 3900: Reconstruction Loss: 0.005855184514075518  
It 4000: Reconstruction Loss: 0.0055256495252251625  
It 4100: Reconstruction Loss: 0.005327051505446434  
It 4200: Reconstruction Loss: 0.005004513077437878  
It 4300: Reconstruction Loss: 0.0063035134226083755  
It 4400: Reconstruction Loss: 0.005292521324008703  
It 4500: Reconstruction Loss: 0.005102543160319328  
It 4600: Reconstruction Loss: 0.0051817456260323524  
Run Epoch 5  
It 4700: Reconstruction Loss: 0.004771128762513399  
It 4800: Reconstruction Loss: 0.005044563673436642  
It 4900: Reconstruction Loss: 0.00499194348230958  
It 5000: Reconstruction Loss: 0.004607779439538717  
It 5100: Reconstruction Loss: 0.005668171215802431  
It 5200: Reconstruction Loss: 0.004626581445336342  
It 5300: Reconstruction Loss: 0.004911825060844421  
It 5400: Reconstruction Loss: 0.005654565989971161  
It 5500: Reconstruction Loss: 0.004666066262871027  
It 5600: Reconstruction Loss: 0.004035145044326782  
Run Epoch 6  
It 5700: Reconstruction Loss: 0.004997056908905506  
It 5800: Reconstruction Loss: 0.00633472204208374



```

It 5900: Reconstruction Loss: 0.004174941219389439
It 6000: Reconstruction Loss: 0.004814298823475838
It 6100: Reconstruction Loss: 0.005388380028307438
It 6200: Reconstruction Loss: 0.004971726331859827
It 6300: Reconstruction Loss: 0.004994853399693966
It 6400: Reconstruction Loss: 0.004561730194836855
It 6500: Reconstruction Loss: 0.0046912203542888165
Run Epoch 7
It 6600: Reconstruction Loss: 0.004468838218599558
It 6700: Reconstruction Loss: 0.004582060966640711
It 6800: Reconstruction Loss: 0.004014728590846062
It 6900: Reconstruction Loss: 0.00393039220944047
It 7000: Reconstruction Loss: 0.00466529093682766
It 7100: Reconstruction Loss: 0.00452636880800128
It 7200: Reconstruction Loss: 0.004711611662060022
It 7300: Reconstruction Loss: 0.0037244956474751234
It 7400: Reconstruction Loss: 0.0045255329459905624
It 7500: Reconstruction Loss: 0.00461358530446887
Run Epoch 8
It 7600: Reconstruction Loss: 0.004474017769098282
It 7700: Reconstruction Loss: 0.0038130246102809906
It 7800: Reconstruction Loss: 0.0045503731817007065
It 7900: Reconstruction Loss: 0.004730247426778078
It 8000: Reconstruction Loss: 0.004492415580898523
It 8100: Reconstruction Loss: 0.004131972324103117
It 8200: Reconstruction Loss: 0.0039808182045817375
It 8300: Reconstruction Loss: 0.004443623591214418
It 8400: Reconstruction Loss: 0.004379660822451115
Run Epoch 9
It 8500: Reconstruction Loss: 0.004476477392017841
It 8600: Reconstruction Loss: 0.004415799863636494
It 8700: Reconstruction Loss: 0.003854422364383936
It 8800: Reconstruction Loss: 0.004195955116301775
It 8900: Reconstruction Loss: 0.003928026184439659
It 9000: Reconstruction Loss: 0.00427329121157527
It 9100: Reconstruction Loss: 0.004124037455767393
It 9200: Reconstruction Loss: 0.003656135406345129
It 9300: Reconstruction Loss: 0.0038930936716496944
~

```

## Verifying reconstructions [Opt]

Now that we trained the auto-encoder we can visualize some of the reconstructions on the test set to verify that it is converged and did not overfit. **Before continuing, make sure that your auto-encoder is able to reconstruct these samples near-perfectly.**


```
In [ ]: # visualize test data reconstructions
def vis_reconstruction(model):
    # download MNIST test set + build Dataset object
    mnist_test = torchvision.datasets.MNIST(root='./data',
                                             train=False,
                                             download=True,
                                             transform=torchvision.transforms.

    mnist_test_iter = iter(mnist_test)
    model.eval() # set model in evaluation mode (eg freeze batchnorm pa
    input_imgs, test_reconstructions = [], []
    for _ in range(5):
        input_img = np.asarray(next(mnist_test_iter)[0])
        reconstruction = model.reconstruct(torch.tensor(input_img[None], device
        input_imgs.append(input_img[0])
        test_reconstructions.append(reconstruction[0, 0].data.cpu().numpy())

    fig = plt.figure(figsize = (20, 50))
    ax1 = plt.subplot(111)
    ax1.imshow(np.concatenate([np.concatenate(input_imgs, axis=1),
                               np.concatenate(test_reconstructions, axis=1)],

    plt.show()

vis_reconstruction(ae_model)
```



## Sampling from the Auto-Encoder [2pt]

To test whether the auto-encoder is useful as a generative model, we can use it like any other generative model: draw embedding samples from a prior distribution and decode them through the decoder network. We will choose a unit Gaussian prior to allow for easy comparison to the VAE later.


```

In [ ]: # we will sample N embeddings, then decode and visualize them
def vis_samples(model):
    ##### TODO #####
    # Sample embeddings from a diagonal unit Gaussian distribution and decode
    # using the model.
    # HINT: The sampled embeddings should have shape [batch_size, nz]. Diagonal
    # Gaussians have mean 0 and a covariance matrix with ones on the diagonal
    # and zeros everywhere else.
    # HINT: If you are unsure whether you sampled the correct distribution, you can
    # sample a large batch and compute the empirical mean and variance using the
    # .mean() and .var() functions.
    # HINT: You can directly use model.decoder() to decode the samples.
    #####
    sampled_embeddings = torch.randn(batch_size, nz).to(device) # sample k
    decoded_samples = model.decoder(sampled_embeddings) # decoder output
    ##### END TODO #####

    fig = plt.figure(figsize = (10, 10))
    ax1 = plt.subplot(111)
    ax1.imshow(torchvision.utils.make_grid(decoded_samples[:16], nrow=4, pad=10,
                                           .data.cpu().numpy().transpose(1, 2, 0), cmap='gray'))
    plt.show()

vis_samples(ae_model)

```



**Inline Question: Describe your observations, why do you think they occur? [2pt]** \ (please limit your answer to <150 words) \ **Answer:** The images generated by the model are not distinct. To be specific the model is trying to generate digits but it is not clearly able to generate a single digit. It seems that two or more digits which are partially generated are overlapping. For example I can see, traces of 3 and 5 overlapping, or may be 4 is converting to 7 or vice versa. Another example is first image in second row has traces of 3 and 6. It can also be noted that the images generated are not in shape and are very blurry.

### 3. Variational Auto-Encoder (VAE)

Variational auto-encoders use a very similar architecture to deterministic auto-encoders, but are inherently stochastic models, i.e. we perform a stochastic sampling operation during the forward pass, leading to different different outputs every time we run the network for the same input. This sampling is required to optimize the VAE objective also known as the evidence lower bound (ELBO):

$$\mathbb{E}_{p(z|x)} [\mathbb{E}_{p(x|z)} [x - \hat{x}]^2] - \mathbb{D}_{\text{KL}}(q(z|x) || p(z))$$

Here,  $\mathbb{D}_{\text{KL}}(q, p)$  denotes the Kullback-Leibler (KL) divergence between the posterior distribution  $q(z|x)$ , i.e. the output of our encoder, and  $p(z)$ , the prior over the embedding variable  $z$ , which we can choose freely.

For simplicity, we will again choose a unit Gaussian prior. The left term is the reconstruction term we already know from training the auto-encoder. When assuming a Gaussian output distribution for both encoder  $q(z|x)$  and decoder  $p(x|z)$  the objective reduces to:

$$\mathcal{L}_{\text{VAE}} = \sum_{x \sim \mathcal{D}} (x - \hat{x})^2 - \beta \cdot \mathbb{D}_{\text{KL}}(\mathcal{N}(\mu_q, \sigma_q) || \mathcal{N}(0, I))$$

Here,  $\hat{x}$  is the reconstruction output of the decoder. In comparison to the auto-encoder objective, the VAE adds a regularizing term between the output of the encoder and a chosen prior distribution, effectively forcing the encoder output to not stray too far from the prior during training. As a result the decoder gets trained with samples that look pretty similar to samples from the prior, which will hopefully allow us to generate better images when using the VAE as a generative model and actually feeding it samples from the prior (as we have done for the AE before).

The coefficient  $\beta$  is a scalar weighting factor that trades off between reconstruction and regularization objective. We will investigate the influence of this factor in our experiments below.

If you need a refresher on VAEs you can check out this tutorial paper: <https://arxiv.org/abs/1606.05908>

## Reparametrization Trick

The sampling procedure inside the VAE's forward pass for obtaining a sample  $z$  from the posterior distribution  $q(z \mid x)$ , when implemented naively, is non-differentiable. However, since  $q(z \mid x)$  is parametrized with a Gaussian function, there is a simple trick to obtain a differentiable sampling operator, known as the *reparametrization trick*.

Instead of directly sampling  $z \sim \mathcal{N}(\mu_q, \sigma_q)$  we can "separate" the network's predictions and the random sampling by computing the sample as:

$$z = \mu_q + \sigma_q * \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

Note that in this equation, the sample  $z$  is computed as a deterministic function of the network's predictions  $\mu_q$  and  $\sigma_q$  and therefore allows to propagate gradients through the sampling procedure.

**Note:** While in the equations above the encoder network parametrizes the standard deviation  $\sigma_q$  of the Gaussian posterior distribution, in practice we usually parametrize the **logarithm of the standard deviation**  $\log \sigma_q$  for numerical stability. Before

## Defining the VAE Model [7pt]

```

In [ ]: def kl_divergence(mu1, log_sigma1, mu2, log_sigma2):
        """Computes KL[p||q] between two Gaussians defined by [mu, log_sigma]."""
        return (log_sigma2 - log_sigma1) + (torch.exp(log_sigma1) ** 2 + (mu1 - mu2) ** 2) / (2 * torch.exp(log_sigma2) ** 2) - 0.5

class VAE(nn.Module):
    def __init__(self, nz, beta=1.0):
        super().__init__()
        self.beta = beta # factor trading off between two loss components
        ##### TODO #####
        # Instantiate Encoder and Decoder.
        # HINT: Remember that the encoder is now parametrizing a Gaussian distribution over the
        # mean and log_sigma, so the dimensionality of the output embedding is
        # double.
        #####
        self.encoder = Encoder(nz * 2)
        self.decoder = Decoder(nz)
        self.nz = nz
        ##### END TODO #####

    def forward(self, x):
        ##### TODO #####
        # Implement the forward pass of the VAE.
        # HINT: Your code should implement the following steps:
        # 1. encode input x, split encoding into mean and log_sigma of the posterior
        # 2. sample z from inferred posterior distribution using reparametrization trick
        # 3. decode the sampled z to obtain the reconstructed image
        #####
        # encode input into posterior distribution q(z | x)
        q = self.encoder(x) # output of encoder (concatenated mean and log_sigma)

        # sample latent variable z with reparametrization
        mu = q[:, : self.nz]
        log_sigma = q[:, self.nz:]

        std = torch.exp(log_sigma)
        eps = torch.randn_like(std)
        z = mu + (eps * std) # batch of sampled embeddings

        # compute reconstruction
        reconstruction = self.decoder(z) # decoder reconstruction from embeddings
        ##### END TODO #####

        return {'q': q,
                'rec': reconstruction}

    def loss(self, x, outputs):
        ##### TODO #####
        # Implement the loss computation of the VAE.
        # HINT: Your code should implement the following steps:
        # 1. compute the image reconstruction loss, similar to AE we use nn.MSELoss()
        # 2. compute the KL divergence loss between the inferred posterior distribution and a unit Gaussian prior; you can use the kl_divergence function above for computing the KL divergence between two Gaussians parametrized by mean and log_sigma
        # HINT: Make sure to compute the KL divergence in the correct order since it is not symmetric, ie. KL(p, q) != KL(q, p)!
        #####
        # compute reconstruction loss
        criterion = nn.MSELoss()
        rec_loss = criterion(x, outputs['rec'])

```

```

rec_loss = criterion(x, outputs[ 'rec' ])

# compute KL divergence loss
q = outputs['q']
mu1 = q[:, 0 : self.nz]
log_sigma1 = q[:, self.nz: ]

# Create a unit gaussian prior
mu2 = torch.zeros_like(mu1)
log_sigma2 = torch.zeros_like(log_sigma1)

kl_loss = kl_divergence(mu1, log_sigma1, mu2, log_sigma2).sum(dim = 1).
##### END TODO #####

# return weighted objective
return rec_loss + self.beta * kl_loss, \
    {'rec_loss': rec_loss, 'kl_loss': kl_loss}

def reconstruct(self, x):
    """Use mean of posterior estimate for visualization reconstruction."""
    ##### TODO #####
    # This function is used for visualizing reconstructions of our VAE model
    # obtain the maximum likelihood estimate we bypass the sampling procedure
    # inferred latent and instead directly use the mean of the inferred posterior
    # HINT: encode the input image and then decode the mean of the posterior
    # the reconstruction.
    #####
    q = self.encoder(x)

    total_len = q.shape
    mid = total_len[1] // 2
    mu1 = q[:, 0 : mid]

    reconstruction = self.decoder(mu1)
    ##### END TODO #####
    return reconstruction

```

## Setting up the VAE Training Loop [4pt]

Let's start training the VAE model! We will first verify our implementation by setting  $\beta = 0$ .

```

In [ ]: learning_rate = 1e-3
        nz = 64

##### TODO #####
# Tune the beta parameter to obtain good VAE training results. However, for
# initial experiments leave beta = 0 in order to verify our implementation.
#####
epochs = 32          # using 5 epochs is sufficient for the first two experi
                     # for the experiment where you tune beta, 20 epochs are

beta = 32e-4
##### END TODO #####

# build VAE model
vae_model = VAE(nz, beta).to(device)    # transfer model to GPU if available
vae_model = vae_model.train()          # set model in train mode (eg batchnorm par

# build optimizer and loss function
##### TODO #####
# Build the optimizer for the vae_model. We will again use the Adam optimizer
# the given learning rate and otherwise default parameters.
#####
opt = torch.optim.Adam(vae_model.parameters(), lr=learning_rate)
##### END TODO #####

train_it = 0
rec_loss, kl_loss = [], []
for ep in range(epochs):
    print("Run Epoch {}".format(ep))
    ##### TODO #####
    # Implement the main training loop for the VAE model.
    # HINT: Your training loop should sample batches from the data loader, run
    #       forward pass of the VAE, compute the loss, perform the backward pass
    #       perform one gradient step with the optimizer.
    # HINT: Don't forget to erase old gradients before performing the backward pass
    # HINT: This time we will use the loss() function of our model for computing
    #       training loss. It outputs the total training loss and a dict containing
    #       the breakdown of reconstruction and KL loss.
    #####
    for sample_img, sample_label in mnist_data_loader:
        # add VAE training loop commands here
        # ...

        # move to device
        sample_img = sample_img.to(device)

        opt.zero_grad()
        outputs = vae_model.forward(sample_img)
        total_loss, losses = vae_model.loss(sample_img, outputs)
        losses['rec_loss'].backward()
        opt.step()

        losses['rec_loss'] = losses['rec_loss'].detach().cpu()
        losses['kl_loss'] = losses['kl_loss'].detach().cpu()

    ##### END TODO #####

    rec_loss.append(losses['rec_loss']); kl_loss.append(losses['kl_loss'])
    if train_it % 100 == 0:
        print("It {}: Total Loss: {}, \t Rec Loss: {}, \t KL Loss: {}".format(
            train_it, total_loss, losses['rec_loss'], losses['kl_loss']))
        train_it += 1

print("Done!")

```



```

print( "Done: " )

# log the loss training curves
fig = plt.figure(figsize = (10, 5))
ax1 = plt.subplot(121)
ax1.plot(rec_loss)
ax1.title.set_text("Reconstruction Loss")
ax2 = plt.subplot(122)
ax2.plot(kl_loss)
ax2.title.set_text("KL Loss")
plt.show()

```

Run Epoch 0

/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:481: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

cpuset\_checked))

It 0: Total Loss: 0.2774864435195923,	Rec Loss: 0.21992328763008118,	KL Loss: 17.988483428955078
It 100: Total Loss: 0.4779055714607239,	Rec Loss: 0.03096154704689	KL Loss: 139.67001342773438
It 200: Total Loss: 0.7319027185440063,	Rec Loss: 0.02000712603330	KL Loss: 222.46737670898438
It 300: Total Loss: 0.8326473236083984,	Rec Loss: 0.01416365895420	KL Loss: 255.77615356445312
It 400: Total Loss: 0.9007357358932495,	Rec Loss: 0.01166407670825	KL Loss: 277.83489990234375
It 500: Total Loss: 0.9605934023857117,	Rec Loss: 0.01222754549235	KL Loss: 296.3643493652344
It 600: Total Loss: 1.0072956085205078,	Rec Loss: 0.01190292369574	KL Loss: 311.0602111816406
It 700: Total Loss: 1.0459448099136353,	Rec Loss: 0.01122049055993	KL Loss: 323.35137939453125
It 800: Total Loss: 1.0774550437927246,	Rec Loss: 0.01151841506361	KL Loss: 333.105224609375
It 900: Total Loss: 1.102566123008728,	Rec Loss: 0.009032628498971462,	KL Loss: 341.7292175292969

Run Epoch 1

It 1000: Total Loss: 1.130315899848938,	Rec Loss: 0.00851168949156	KL Loss: 350.5638427734375
It 1100: Total Loss: 1.151742935180664,	Rec Loss: 0.00939492974430	KL Loss: 356.9837646484375
It 1200: Total Loss: 1.1788643598556519,	Rec Loss: 0.00796042289584	KL Loss: 365.907470703125
It 1300: Total Loss: 1.2035024166107178,	Rec Loss: 0.00856794696301	KL Loss: 373.41705322265625
It 1400: Total Loss: 1.2265050411224365,	Rec Loss: 0.00832919869571	KL Loss: 380.6799621582031
It 1500: Total Loss: 1.2449636459350586,	Rec Loss: 0.00804529525339	KL Loss: 386.5369873046875
It 1600: Total Loss: 1.2702120542526245,	Rec Loss: 0.00830254796892	KL Loss: 394.34674072265625
It 1700: Total Loss: 1.2851167917251587,	Rec Loss: 0.00677153049036	KL Loss: 399.48291015625
It 1800: Total Loss: 1.306889295578003,	Rec Loss: 0.00745501695200	KL Loss: 406.0732421875

Run Epoch 2

It 1900: Total Loss: 1.323966383934021,	Rec Loss: 0.00702805723994	KL Loss: 411.543212890625
It 2000: Total Loss: 1.3477437496185303,	Rec Loss: 0.00730635784566	KL Loss: 418.8866882324219

It 2100: Total Loss: 1.3664422035217285,	Rec Loss: 0.00658883480355
14355, KL Loss: 424.9541931152344	
It 2200: Total Loss: 1.38427734375,	Rec Loss: 0.006209683604538441,
KL Loss: 430.6461486816406	
It 2300: Total Loss: 1.4032166004180908,	Rec Loss: 0.00718410266563
2963, KL Loss: 436.2601623535156	
It 2400: Total Loss: 1.4290953874588013,	Rec Loss: 0.00711045693606
1382, KL Loss: 444.37030029296875	
It 2500: Total Loss: 1.4324791431427002,	Rec Loss: 0.005608934444344
4014, KL Loss: 445.8969421386719	
It 2600: Total Loss: 1.4486968517303467,	Rec Loss: 0.00665817270055
41325, KL Loss: 450.6370849609375	
It 2700: Total Loss: 1.4847843647003174,	Rec Loss: 0.00676075648516
41655, KL Loss: 461.8824157714844	
It 2800: Total Loss: 1.5002268552780151,	Rec Loss: 0.00663118204101
9201, KL Loss: 466.7486877441406	
Run Epoch 3	
It 2900: Total Loss: 1.512379765510559,	Rec Loss: 0.00589347071945
6673, KL Loss: 470.7769775390625	
It 3000: Total Loss: 1.5423152446746826,	Rec Loss: 0.00712254922837
019, KL Loss: 479.74774169921875	
It 3100: Total Loss: 1.5565309524536133,	Rec Loss: 0.00572328083217
144, KL Loss: 484.62744140625	
It 3200: Total Loss: 1.577323079109192,	Rec Loss: 0.00692464830353
8561, KL Loss: 490.74951171875	
It 3300: Total Loss: 1.5972665548324585,	Rec Loss: 0.00765071902424
0971, KL Loss: 496.75494384765625	
It 3400: Total Loss: 1.6176013946533203,	Rec Loss: 0.00693783676251
7691, KL Loss: 503.3323669433594	
It 3500: Total Loss: 1.613166093826294,	Rec Loss: 0.00564650585874
9151, KL Loss: 502.34991455078125	
It 3600: Total Loss: 1.6424843072891235,	Rec Loss: 0.00537592638283
968, KL Loss: 511.59637451171875	
It 3700: Total Loss: 1.6607036590576172,	Rec Loss: 0.00585298426449
2989, KL Loss: 517.140869140625	
Run Epoch 4	
It 3800: Total Loss: 1.668525218963623,	Rec Loss: 0.00663072336465
1203, KL Loss: 519.342041015625	
It 3900: Total Loss: 1.6916613578796387,	Rec Loss: 0.00615688459947
7053, KL Loss: 526.7201538085938	
It 4000: Total Loss: 1.714827299118042,	Rec Loss: 0.00542154768481
8506, KL Loss: 534.1893310546875	
It 4100: Total Loss: 1.7321176528930664,	Rec Loss: 0.00566127803176
6415, KL Loss: 539.5176391601562	
It 4200: Total Loss: 1.7553743124008179,	Rec Loss: 0.00607981579378
2473, KL Loss: 546.654541015625	
It 4300: Total Loss: 1.7606512308120728,	Rec Loss: 0.00482209911569
953, KL Loss: 548.6965942382812	
It 4400: Total Loss: 1.7789299488067627,	Rec Loss: 0.00541949039325
1181, KL Loss: 554.2220458984375	
It 4500: Total Loss: 1.786085844039917,	Rec Loss: 0.00494317244738
3404, KL Loss: 556.6071166992188	
It 4600: Total Loss: 1.8270859718322754,	Rec Loss: 0.00565514015033
8411, KL Loss: 569.1971435546875	
Run Epoch 5	
It 4700: Total Loss: 1.843637466430664,	Rec Loss: 0.00556664774194
3598, KL Loss: 574.3971557617188	
It 4800: Total Loss: 1.858289122581482,	Rec Loss: 0.00632513640448
451, KL Loss: 578.73876953125	
It 4900: Total Loss: 1.8836029767990112,	Rec Loss: 0.00499674677848
8159, KL Loss: 587.064453125	
It 5000: Total Loss: 1.902522087097168,	Rec Loss: 0.00481174793094
3966, KL Loss: 593.0344848632812	

It 5100: Total Loss: 1.9192194938659668, 2078, KL Loss: 598.237548828125	Rec Loss: 0.00485937297344
It 5200: Total Loss: 1.936258316040039, 2156, KL Loss: 603.543701171875	Rec Loss: 0.00491855246946
It 5300: Total Loss: 1.9503813982009888, 6669, KL Loss: 608.0143432617188	Rec Loss: 0.00473559694364
It 5400: Total Loss: 1.975095510482788, 0565, KL Loss: 615.4722900390625	Rec Loss: 0.00558425579220
It 5500: Total Loss: 1.9986612796783447, 5066, KL Loss: 622.817138671875	Rec Loss: 0.00564643926918
It 5600: Total Loss: 2.0086140632629395, 4469, KL Loss: 626.0599365234375	Rec Loss: 0.00522227492183
Run Epoch 6	
It 5700: Total Loss: 2.02952241897583, KL Loss: 632.8350830078125	Rec Loss: 0.0044502634555101395,
It 5800: Total Loss: 2.0493009090423584, 8647, KL Loss: 638.9962768554688	Rec Loss: 0.00451282458379
It 5900: Total Loss: 2.0689947605133057, 245, KL Loss: 645.013671875	Rec Loss: 0.00495099276304
It 6000: Total Loss: 2.094130039215088, 4936, KL Loss: 652.94873046875	Rec Loss: 0.00469426112249
It 6100: Total Loss: 2.114453077316284, 943, KL Loss: 658.9991455078125	Rec Loss: 0.00565573014318
It 6200: Total Loss: 2.1350691318511963, 2697, KL Loss: 665.662109375	Rec Loss: 0.00495060579851
It 6300: Total Loss: 2.1467838287353516, 8203, KL Loss: 669.2481079101562	Rec Loss: 0.00519000971689
It 6400: Total Loss: 2.187685489654541, 231, KL Loss: 681.8792724609375	Rec Loss: 0.00567197659984
It 6500: Total Loss: 2.180095911026001, 6588, KL Loss: 679.6921997070312	Rec Loss: 0.00508099375292
Run Epoch 7	
It 6600: Total Loss: 2.2090466022491455, 5183, KL Loss: 688.9298706054688	Rec Loss: 0.00447110831737
It 6700: Total Loss: 2.2277493476867676, 0031, KL Loss: 694.8063354492188	Rec Loss: 0.00436911592260
It 6800: Total Loss: 2.2506725788116455, 0575, KL Loss: 701.8640747070312	Rec Loss: 0.00470764003694
It 6900: Total Loss: 2.2832400798797607, 5181, KL Loss: 711.9765014648438	Rec Loss: 0.00491534499451
It 7000: Total Loss: 2.2968590259552, KL Loss: 716.1087646484375	Rec Loss: 0.0053110974840819836,
It 7100: Total Loss: 2.3148977756500244, 1236, KL Loss: 721.6947631835938	Rec Loss: 0.00547448825091
It 7200: Total Loss: 2.3346030712127686, 7329, KL Loss: 727.9039916992188	Rec Loss: 0.00531032169237
It 7300: Total Loss: 2.373723268508911, 49795, KL Loss: 740.3370361328125	Rec Loss: 0.00464492803439
It 7400: Total Loss: 2.3963327407836914, 5468, KL Loss: 747.3071899414062	Rec Loss: 0.00494982907548
It 7500: Total Loss: 2.4058775901794434, 2774, KL Loss: 750.3427124023438	Rec Loss: 0.00478109065443
Run Epoch 8	
It 7600: Total Loss: 2.436552047729492, 2226, KL Loss: 759.9234619140625	Rec Loss: 0.00479691987857
It 7700: Total Loss: 2.4529664516448975, 62965, KL Loss: 765.0106811523438	Rec Loss: 0.00493241194635
It 7800: Total Loss: 2.481835126876831, 8197, KL Loss: 774.0535278320312	Rec Loss: 0.00486395834013
It 7900: Total Loss: 2.5118844509124756, 5027, KL Loss: 783.5551147460938	Rec Loss: 0.00450830161571
It 8000: Total Loss: 2.5399680137634277, 2588, KL Loss: 792.0643310546875	Rec Loss: 0.00536237331107

It 8100: Total Loss: 2.552361249923706,	Rec Loss: 0.00502146454527
9741, KL Loss: 796.043701171875	
It 8200: Total Loss: 2.5769786834716797,	Rec Loss: 0.00447487132623
7917, KL Loss: 803.907470703125	
It 8300: Total Loss: 2.589060068130493,	Rec Loss: 0.00425026845186
9488, KL Loss: 807.7530517578125	
It 8400: Total Loss: 2.6047215461730957,	Rec Loss: 0.00404173368588
0899, KL Loss: 812.7124633789062	
Run Epoch 9	
It 8500: Total Loss: 2.637462854385376,	Rec Loss: 0.00411361269652
8435, KL Loss: 822.921630859375	
It 8600: Total Loss: 2.6765949726104736,	Rec Loss: 0.00460720062255
8594, KL Loss: 834.9962158203125	
It 8700: Total Loss: 2.699052095413208,	Rec Loss: 0.00471660494804
3823, KL Loss: 841.9798583984375	
It 8800: Total Loss: 2.7238516807556152,	Rec Loss: 0.00470954738557
3387, KL Loss: 849.73193359375	
It 8900: Total Loss: 2.749663829803467,	Rec Loss: 0.00418854970484
972, KL Loss: 857.9610595703125	
It 9000: Total Loss: 2.7636239528656006,	Rec Loss: 0.00443407986313
1046, KL Loss: 862.246826171875	
It 9100: Total Loss: 2.79076886177063,	Rec Loss: 0.004247513134032488,
KL Loss: 870.7879638671875	
It 9200: Total Loss: 2.802940845489502,	Rec Loss: 0.00416485872119
66515, KL Loss: 874.6174926757812	
It 9300: Total Loss: 2.832237720489502,	Rec Loss: 0.00455129146575
9277, KL Loss: 883.652099609375	
Run Epoch 10	
It 9400: Total Loss: 2.8359737396240234,	Rec Loss: 0.00471646245568
9907, KL Loss: 884.7679443359375	
It 9500: Total Loss: 2.893840789794922,	Rec Loss: 0.00469406228512
5256, KL Loss: 902.8583984375	
It 9600: Total Loss: 2.921339273452759,	Rec Loss: 0.00460641644895
07675, KL Loss: 911.47900390625	
It 9700: Total Loss: 2.93939471244812,	Rec Loss: 0.004112901631742716,
KL Loss: 917.2755737304688	
It 9800: Total Loss: 2.965461254119873,	Rec Loss: 0.00404720706865
1915, KL Loss: 925.44189453125	
It 9900: Total Loss: 2.9857583045959473,	Rec Loss: 0.00437527196481
8239, KL Loss: 931.6822509765625	
It 10000: Total Loss: 3.025643825531006,	Rec Loss: 0.00448724534362
55455, KL Loss: 944.1114501953125	
It 10100: Total Loss: 3.027818202972412,	Rec Loss: 0.00398918474093
0796, KL Loss: 944.9465942382812	
It 10200: Total Loss: 3.0687053203582764,	Rec Loss: 0.00510968640446
6629, KL Loss: 957.3736572265625	
It 10300: Total Loss: 3.0727362632751465,	Rec Loss: 0.00390825280919
6711, KL Loss: 959.0087890625	
Run Epoch 11	
It 10400: Total Loss: 3.0769262313842773,	Rec Loss: 0.00363808963447
8092, KL Loss: 960.402587890625	
It 10500: Total Loss: 3.110595464706421,	Rec Loss: 0.00417428417131
3047, KL Loss: 970.7566528320312	
It 10600: Total Loss: 3.137021541595459,	Rec Loss: 0.00387955061160
028, KL Loss: 979.1068725585938	
It 10700: Total Loss: 3.1840057373046875,	Rec Loss: 0.00458821700885
8919, KL Loss: 993.5680541992188	
It 10800: Total Loss: 3.1893532276153564,	Rec Loss: 0.00411293981596
8275, KL Loss: 995.3876342773438	
It 10900: Total Loss: 3.2263028621673584,	Rec Loss: 0.00430353032425
046, KL Loss: 1006.8748168945312	
It 11000: Total Loss: 3.259866714477539,	Rec Loss: 0.00410383008420
4674, KL Loss: 1017.4259033203125	

It 11100: Total Loss: 3.28562068939209, 56445, KL Loss: 1025.378662109375	Rec Loss: 0.00440895557403
It 11200: Total Loss: 3.315359115600586, 1428, KL Loss: 1034.673828125	Rec Loss: 0.00440288055688
Run Epoch 12	
It 11300: Total Loss: 3.3569817543029785, 5524, KL Loss: 1047.812255859375	Rec Loss: 0.00398242566734
It 11400: Total Loss: 3.338604211807251, 8389, KL Loss: 1042.0205078125	Rec Loss: 0.00413873139768
It 11500: Total Loss: 3.363919734954834, 8626, KL Loss: 1049.954345703125	Rec Loss: 0.00406595086678
It 11600: Total Loss: 3.395744562149048, 8931, KL Loss: 1060.047607421875	Rec Loss: 0.00359237054362
It 11700: Total Loss: 3.398911952972412, 9459, KL Loss: 1060.952392578125	Rec Loss: 0.00386428157798
It 11800: Total Loss: 3.428546905517578, 9337, KL Loss: 1070.1241455078125	Rec Loss: 0.00414969585835
It 11900: Total Loss: 3.4480550289154053, 3049, KL Loss: 1076.158447265625	Rec Loss: 0.00434812204912
It 12000: Total Loss: 3.465754508972168, 4025, KL Loss: 1081.6949462890625	Rec Loss: 0.00433083157986
It 12100: Total Loss: 3.4740633964538574, 4215, KL Loss: 1084.4122314453125	Rec Loss: 0.00394451199099
Run Epoch 13	
It 12200: Total Loss: 3.522129535675049, 4577, KL Loss: 1099.48095703125	Rec Loss: 0.00379054783843
It 12300: Total Loss: 3.5470144748687744, 4092, KL Loss: 1107.082763671875	Rec Loss: 0.00434980448335
It 12400: Total Loss: 3.564781427383423, 6185, KL Loss: 1112.6875	Rec Loss: 0.00418167188763
It 12500: Total Loss: 3.568160057067871, 78133, KL Loss: 1113.8876953125	Rec Loss: 0.00371964578516
It 12600: Total Loss: 3.611981153488159, 4667, KL Loss: 1127.618896484375	Rec Loss: 0.00360085163265
It 12700: Total Loss: 3.632220506668091, 63982, KL Loss: 1133.869140625	Rec Loss: 0.00383919430896
It 12800: Total Loss: 3.644129991531372, 10595, KL Loss: 1137.6458740234375	Rec Loss: 0.00366319133900
It 12900: Total Loss: 3.6695828437805176, 90216, KL Loss: 1145.54931640625	Rec Loss: 0.00382528663612
It 13000: Total Loss: 3.68013596534729, 0489, KL Loss: 1148.836181640625	Rec Loss: 0.00386013649404
It 13100: Total Loss: 3.7122390270233154, 33847, KL Loss: 1158.85546875	Rec Loss: 0.00390179106034
Run Epoch 14	
It 13200: Total Loss: 3.7361929416656494, 88243, KL Loss: 1166.501953125	Rec Loss: 0.00338671100325
It 13300: Total Loss: 3.7613213062286377, 5713, KL Loss: 1174.3067626953125	Rec Loss: 0.00353970471769
It 13400: Total Loss: 3.7764878273010254, 7135, KL Loss: 1178.713134765625	Rec Loss: 0.00460598105564
It 13500: Total Loss: 3.7983922958374023, 7026, KL Loss: 1185.72998046875	Rec Loss: 0.00405638245865
It 13600: Total Loss: 3.820723295211792, 1914, KL Loss: 1192.755126953125	Rec Loss: 0.00390685116872
It 13700: Total Loss: 3.8380980491638184, 5416, KL Loss: 1198.22607421875	Rec Loss: 0.00377457402646
It 13800: Total Loss: 3.876309871673584, 89303, KL Loss: 1210.3243408203125	Rec Loss: 0.00327199837192
It 13900: Total Loss: 3.8860723972320557, 3448, KL Loss: 1213.0723876953125	Rec Loss: 0.00424073822796
It 14000: Total Loss: 3.918910264968872, 7484, KL Loss: 1223.4705810546875	Rec Loss: 0.00380454794503

Run Epoch 15

It 14100: Total Loss: 3.926501750946045, 0014, KL Loss: 1225.5611572265625	Rec Loss: 0.00470618205145
It 14200: Total Loss: 3.955382823944092, 40576, KL Loss: 1234.908935546875	Rec Loss: 0.00367433740757
It 14300: Total Loss: 3.9658830165863037, 2724, KL Loss: 1238.245361328125	Rec Loss: 0.00349806458689
It 14400: Total Loss: 4.005494594573975, 11963, KL Loss: 1250.530029296875	Rec Loss: 0.00379849388264
It 14500: Total Loss: 4.01389741897583, 7983, KL Loss: 1253.077880859375	Rec Loss: 0.00404831254854
It 14600: Total Loss: 3.988095760345459, 4069, KL Loss: 1245.29296875	Rec Loss: 0.00315840169787
It 14700: Total Loss: 4.025870323181152, 7975, KL Loss: 1256.8602294921875	Rec Loss: 0.00391772342845
It 14800: Total Loss: 4.045266151428223, 32964, KL Loss: 1262.94677734375	Rec Loss: 0.00383653445169
It 14900: Total Loss: 4.0716352462768555, 4881, KL Loss: 1271.213623046875	Rec Loss: 0.00375193660147
It 15000: Total Loss: 4.083459377288818, 1492, KL Loss: 1275.072265625	Rec Loss: 0.00322805601172

Run Epoch 16

It 15100: Total Loss: 4.131476402282715, 93145, KL Loss: 1289.774169921875	Rec Loss: 0.00419882638379
It 15200: Total Loss: 4.128103733062744, 3776, KL Loss: 1288.923095703125	Rec Loss: 0.00355011830106
It 15300: Total Loss: 4.160480499267578, 2481, KL Loss: 1299.027587890625	Rec Loss: 0.00359271909110
It 15400: Total Loss: 4.171641826629639, 24563, KL Loss: 1302.62255859375	Rec Loss: 0.00324987969361
It 15500: Total Loss: 4.200206279754639, 6252, KL Loss: 1311.28564453125	Rec Loss: 0.00409222627058
It 15600: Total Loss: 4.196268081665039, 23286, KL Loss: 1310.1473388671875	Rec Loss: 0.00379667477682
It 15700: Total Loss: 4.225154399871826, 4522, KL Loss: 1319.12744140625	Rec Loss: 0.00394685612991
It 15800: Total Loss: 4.241907119750977, 38525, KL Loss: 1324.4508056640625	Rec Loss: 0.00366440415382
It 15900: Total Loss: 4.257602214813232, 8625, KL Loss: 1329.1895751953125	Rec Loss: 0.00419574417173

Run Epoch 17

It 16000: Total Loss: 4.241065979003906, 55365, KL Loss: 1324.3348388671875	Rec Loss: 0.00319479615427
It 16100: Total Loss: 4.294215679168701, 70677, KL Loss: 1340.857421875	Rec Loss: 0.00347163435071
It 16200: Total Loss: 4.313192844390869, 0836, KL Loss: 1346.739990234375	Rec Loss: 0.00362501037307
It 16300: Total Loss: 4.309237003326416, 85207, KL Loss: 1345.5989990234375	Rec Loss: 0.00331999291665
It 16400: Total Loss: 4.360329627990723, 4259, KL Loss: 1361.472412109375	Rec Loss: 0.00361820380203
It 16500: Total Loss: 4.3884077072143555, 05236, KL Loss: 1370.35498046875	Rec Loss: 0.00327215087600
It 16600: Total Loss: 4.380279064178467, 12627, KL Loss: 1367.8443603515625	Rec Loss: 0.00317736994475
It 16700: Total Loss: 4.441582202911377, 99553, KL Loss: 1386.822509765625	Rec Loss: 0.00375045533291
It 16800: Total Loss: 4.444010257720947, 8767, KL Loss: 1387.46337890625	Rec Loss: 0.00412761187180

Run Epoch 18

It 16900: Total Loss: 4.457760334014893, 1345, KL Loss: 1391.92578125	Rec Loss: 0.00359787698835
It 17000: Total Loss: 4.467045307159424,	Rec Loss: 0.00335479271598

16027, KL Loss: 1394.9033203125	
It 17100: Total Loss: 4.48316764831543, 4484, KL Loss: 1399.882568359375	Rec Loss: 0.00354333245195
It 17200: Total Loss: 4.505029678344727, 42056, KL Loss: 1406.664794921875	Rec Loss: 0.00370266474783
It 17300: Total Loss: 4.521445274353027, 92303, KL Loss: 1411.95166015625	Rec Loss: 0.00320015475153
It 17400: Total Loss: 4.561147212982178, 33723, KL Loss: 1424.275146484375	Rec Loss: 0.00346717284992
It 17500: Total Loss: 4.568474292755127, 68575, KL Loss: 1426.418701171875	Rec Loss: 0.00393438804894
It 17600: Total Loss: 4.566519737243652, 83325, KL Loss: 1426.067138671875	Rec Loss: 0.00310497893951
It 17700: Total Loss: 4.6302008628845215, 93755, KL Loss: 1445.7607421875	Rec Loss: 0.00376658840104
It 17800: Total Loss: 4.61329984664917, 097, KL Loss: 1440.5830078125	Rec Loss: 0.00343401194550
Run Epoch 19	
It 17900: Total Loss: 4.634116172790527, 24288, KL Loss: 1447.132080078125	Rec Loss: 0.00329333101399
It 18000: Total Loss: 4.658857822418213, 45276, KL Loss: 1454.7125244140625	Rec Loss: 0.00377780036069
It 18100: Total Loss: 4.636510372161865, 91183, KL Loss: 1447.870361328125	Rec Loss: 0.00332528166472
It 18200: Total Loss: 4.667378902435303, 4371, KL Loss: 1457.598388671875	Rec Loss: 0.00306428619660
It 18300: Total Loss: 4.7315449714660645, 0335, KL Loss: 1477.6324462890625	Rec Loss: 0.00312126125209
It 18400: Total Loss: 4.714032173156738, 51475, KL Loss: 1472.056884765625	Rec Loss: 0.00345015758648
It 18500: Total Loss: 4.771265029907227, 8148, KL Loss: 1489.733154296875	Rec Loss: 0.00411878898739
It 18600: Total Loss: 4.783430099487305, 9465, KL Loss: 1493.78564453125	Rec Loss: 0.00331601058132
It 18700: Total Loss: 4.754288673400879, 18213, KL Loss: 1484.6700439453125	Rec Loss: 0.00334454490803
Run Epoch 20	
It 18800: Total Loss: 4.758532524108887, 49557, KL Loss: 1486.0897216796875	Rec Loss: 0.00304546183906
It 18900: Total Loss: 4.797486782073975, 30905, KL Loss: 1498.0709228515625	Rec Loss: 0.00365958921611
It 19000: Total Loss: 4.823488712310791, 1037, KL Loss: 1506.123046875	Rec Loss: 0.00389542733319
It 19100: Total Loss: 4.8221354484558105, 12684, KL Loss: 1505.838134765625	Rec Loss: 0.00345384585671
It 19200: Total Loss: 4.8300299644470215, 00936, KL Loss: 1508.230712890625	Rec Loss: 0.00369184021838
It 19300: Total Loss: 4.846795082092285, 6221, KL Loss: 1513.5048828125	Rec Loss: 0.00357978488318
It 19400: Total Loss: 4.872220516204834, 81455, KL Loss: 1521.6016845703125	Rec Loss: 0.00309509271755
It 19500: Total Loss: 4.8929243087768555, 76795, KL Loss: 1527.988525390625	Rec Loss: 0.00336132780648
It 19600: Total Loss: 4.938602924346924, 55367, KL Loss: 1542.302734375	Rec Loss: 0.00323445652611
Run Epoch 21	
It 19700: Total Loss: 4.973816394805908, 5743, KL Loss: 1553.205078125	Rec Loss: 0.00356012815609
It 19800: Total Loss: 4.950754642486572, 16276, KL Loss: 1546.038818359375	Rec Loss: 0.00343057210557
It 19900: Total Loss: 4.954117774963379, 2426, KL Loss: 1547.04443359375	Rec Loss: 0.00357583700679
It 20000: Total Loss: 4.984787940979004,	Rec Loss: 0.00306142703630

03016,	KL Loss: 1556.7896728515625	
It 20100: Total Loss: 5.045109748840332,		Rec Loss: 0.00375230936333
5371,	KL Loss: 1575.42431640625	
It 20200: Total Loss: 5.040319442749023,		Rec Loss: 0.00305088004097
3425,	KL Loss: 1574.146484375	
It 20300: Total Loss: 5.0455169677734375,		Rec Loss: 0.00350734149105
8469,	KL Loss: 1575.6280517578125	
It 20400: Total Loss: 5.068269729614258,		Rec Loss: 0.00281632854603
23095,	KL Loss: 1582.9542236328125	
It 20500: Total Loss: 5.051268577575684,		Rec Loss: 0.00332032400183
37965,	KL Loss: 1577.48388671875	
It 20600: Total Loss: 5.095432281494141,		Rec Loss: 0.00339265260845
42274,	KL Loss: 1591.262451171875	
Run Epoch 22		
It 20700: Total Loss: 5.107320308685303,		Rec Loss: 0.00321612157858
90818,	KL Loss: 1595.0325927734375	
It 20800: Total Loss: 5.112602710723877,		Rec Loss: 0.00294342194683
8498,	KL Loss: 1596.7685546875	
It 20900: Total Loss: 5.127891540527344,		Rec Loss: 0.00366360181942
58213,	KL Loss: 1601.3212890625	
It 21000: Total Loss: 5.144627571105957,		Rec Loss: 0.00329956132918
59627,	KL Loss: 1606.6650390625	
It 21100: Total Loss: 5.138282299041748,		Rec Loss: 0.00347360759042
20343,	KL Loss: 1604.627685546875	
It 21200: Total Loss: 5.182568073272705,		Rec Loss: 0.00297306175343
6923,	KL Loss: 1618.62353515625	
It 21300: Total Loss: 5.208220481872559,		Rec Loss: 0.00352404336445
03355,	KL Loss: 1626.4677734375	
It 21400: Total Loss: 5.195047855377197,		Rec Loss: 0.00312364823184
90744,	KL Loss: 1622.476318359375	
It 21500: Total Loss: 5.180099010467529,		Rec Loss: 0.00319016794674
0985,	KL Loss: 1617.7840576171875	
Run Epoch 23		
It 21600: Total Loss: 5.222617149353027,		Rec Loss: 0.00316396402195
096,	KL Loss: 1631.0792236328125	
It 21700: Total Loss: 5.205334186553955,		Rec Loss: 0.00311193917877
97213,	KL Loss: 1625.694580078125	
It 21800: Total Loss: 5.28769063949585,		Rec Loss: 0.00329313962720
33453,	KL Loss: 1651.374267578125	
It 21900: Total Loss: 5.246123790740967,		Rec Loss: 0.00366227468475
69942,	KL Loss: 1638.269287109375	
It 22000: Total Loss: 5.244804859161377,		Rec Loss: 0.00291131762787
6997,	KL Loss: 1638.091796875	
It 22100: Total Loss: 5.304995059967041,		Rec Loss: 0.00339138205163
18083,	KL Loss: 1656.751220703125	
It 22200: Total Loss: 5.3204498291015625,		Rec Loss: 0.00332047441042
9597,	KL Loss: 1661.6029052734375	
It 22300: Total Loss: 5.310042381286621,		Rec Loss: 0.00308415642939
50796,	KL Loss: 1658.4244384765625	
It 22400: Total Loss: 5.332967281341553,		Rec Loss: 0.00300709996372
4613,	KL Loss: 1665.6126708984375	
It 22500: Total Loss: 5.369832992553711,		Rec Loss: 0.00318542239256
20317,	KL Loss: 1677.077392578125	
Run Epoch 24		
It 22600: Total Loss: 5.353749752044678,		Rec Loss: 0.00263645336963
23633,	KL Loss: 1672.222900390625	
It 22700: Total Loss: 5.3684258460998535,		Rec Loss: 0.00329772662371
397,	KL Loss: 1676.6025390625	
It 22800: Total Loss: 5.376986026763916,		Rec Loss: 0.00296518625691
5331,	KL Loss: 1679.381591796875	
It 22900: Total Loss: 5.406606674194336,		Rec Loss: 0.00307878735475
24214,	KL Loss: 1688.6024169921875	
It 23000: Total Loss: 5.449256896972656,		Rec Loss: 0.00325480941683



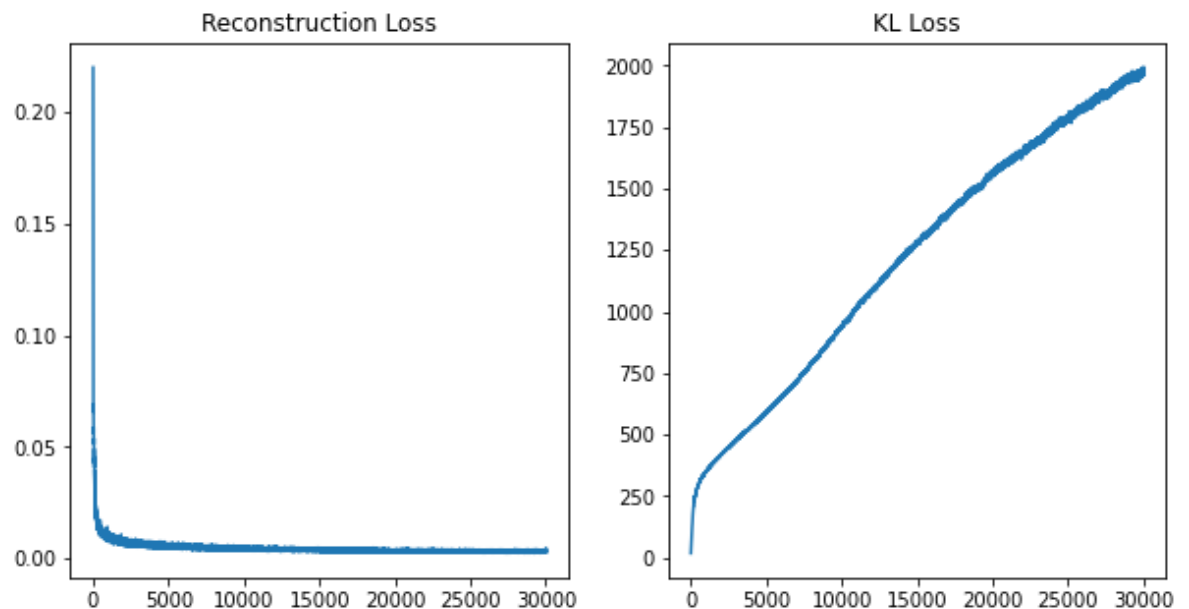
05397, KL Loss: 1701.8756103515625	
It 23100: Total Loss: 5.427193641662598,	Rec Loss: 0.00277926912531
25668, KL Loss: 1695.12939453125	
It 23200: Total Loss: 5.422658920288086,	Rec Loss: 0.00322388252243
3996, KL Loss: 1693.573486328125	
It 23300: Total Loss: 5.464754581451416,	Rec Loss: 0.00295797525905
0727, KL Loss: 1706.8115234375	
It 23400: Total Loss: 5.454076290130615,	Rec Loss: 0.00294203218072
6528, KL Loss: 1703.4794921875	
Run Epoch 25	
It 23500: Total Loss: 5.493338108062744,	Rec Loss: 0.00315974792465
5676, KL Loss: 1715.680908203125	
It 23600: Total Loss: 5.506298065185547,	Rec Loss: 0.00301530095748
60334, KL Loss: 1719.77587890625	
It 23700: Total Loss: 5.51902961730957,	Rec Loss: 0.00297467363998
29388, KL Loss: 1723.7672119140625	
It 23800: Total Loss: 5.533520221710205,	Rec Loss: 0.00300449272617
6977, KL Loss: 1728.2861328125	
It 23900: Total Loss: 5.537171840667725,	Rec Loss: 0.00313666020520
03145, KL Loss: 1729.385986328125	
It 24000: Total Loss: 5.57315731048584,	Rec Loss: 0.00295172957703
4712, KL Loss: 1740.6893310546875	
It 24100: Total Loss: 5.6025214195251465,	Rec Loss: 0.00330907804891
4671, KL Loss: 1749.75390625	
It 24200: Total Loss: 5.6038079261779785,	Rec Loss: 0.00301627232693
13574, KL Loss: 1750.2474365234375	
It 24300: Total Loss: 5.628164291381836,	Rec Loss: 0.00320328609086
57312, KL Loss: 1757.80029296875	
Run Epoch 26	
It 24400: Total Loss: 5.657005310058594,	Rec Loss: 0.00267876894213
2592, KL Loss: 1766.97705078125	
It 24500: Total Loss: 5.696140766143799,	Rec Loss: 0.00300815701484
68018, KL Loss: 1779.1038818359375	
It 24600: Total Loss: 5.697236061096191,	Rec Loss: 0.00340454280376
43433, KL Loss: 1779.3223876953125	
It 24700: Total Loss: 5.65086030960083,	Rec Loss: 0.00317359459586
4415, KL Loss: 1764.902099609375	
It 24800: Total Loss: 5.666318416595459,	Rec Loss: 0.00311747193336
4868, KL Loss: 1769.7503662109375	
It 24900: Total Loss: 5.6986517906188965,	Rec Loss: 0.00320714921690
52362, KL Loss: 1779.826416015625	
It 25000: Total Loss: 5.73121976852417,	Rec Loss: 0.00319123500958
08506, KL Loss: 1790.0089111328125	
It 25100: Total Loss: 5.746094703674316,	Rec Loss: 0.00329975970089
43558, KL Loss: 1794.62353515625	
It 25200: Total Loss: 5.728720664978027,	Rec Loss: 0.00322350906208
15754, KL Loss: 1789.2178955078125	
It 25300: Total Loss: 5.741361141204834,	Rec Loss: 0.00297206337563
69352, KL Loss: 1793.24658203125	
Run Epoch 27	
It 25400: Total Loss: 5.778127670288086,	Rec Loss: 0.00284880748949
94497, KL Loss: 1804.7747802734375	
It 25500: Total Loss: 5.769482135772705,	Rec Loss: 0.00323353148996
83, KL Loss: 1801.9527587890625	
It 25600: Total Loss: 5.798570156097412,	Rec Loss: 0.00299098505638
5398, KL Loss: 1811.1185302734375	
It 25700: Total Loss: 5.844773769378662,	Rec Loss: 0.00312819890677
92892, KL Loss: 1825.514404296875	
It 25800: Total Loss: 5.849793910980225,	Rec Loss: 0.00311080995015
8, KL Loss: 1827.0885009765625	
It 25900: Total Loss: 5.842171669006348,	Rec Loss: 0.00352330692112
44583, KL Loss: 1824.57763671875	
It 26000: Total Loss: 5.879512310028076,	Rec Loss: 0.00305236107669

7707,	KL Loss: 1836.393798828125	
It 26100: Total Loss: 5.864948272705078,		Rec Loss: 0.00284939259290
6952,	KL Loss: 1831.9058837890625	
It 26200: Total Loss: 5.877519607543945,		Rec Loss: 0.00302530359476
80473,	KL Loss: 1835.7794189453125	
Run Epoch 28		
It 26300: Total Loss: 5.892045974731445,		Rec Loss: 0.00335674569942
05713,	KL Loss: 1840.21533203125	
It 26400: Total Loss: 5.90395975112915,		Rec Loss: 0.00322253606282
1746,	KL Loss: 1843.98046875	
It 26500: Total Loss: 5.902001857757568,		Rec Loss: 0.00301938992924
98827,	KL Loss: 1843.43212890625	
It 26600: Total Loss: 5.908114433288574,		Rec Loss: 0.00308642443269
4912,	KL Loss: 1845.3212890625	
It 26700: Total Loss: 5.970120429992676,		Rec Loss: 0.00384446303360
16417,	KL Loss: 1864.4613037109375	
It 26800: Total Loss: 6.005579471588135,		Rec Loss: 0.00330705242231
48823,	KL Loss: 1875.710205078125	
It 26900: Total Loss: 5.960150718688965,		Rec Loss: 0.00314529985189
43787,	KL Loss: 1861.564208984375	
It 27000: Total Loss: 5.999385356903076,		Rec Loss: 0.00285297073423
86246,	KL Loss: 1873.91650390625	
It 27100: Total Loss: 5.988559246063232,		Rec Loss: 0.00302097154781
22234,	KL Loss: 1870.4808349609375	
It 27200: Total Loss: 6.058098793029785,		Rec Loss: 0.00286664580926
2991,	KL Loss: 1892.260009765625	
Run Epoch 29		
It 27300: Total Loss: 6.001765251159668,		Rec Loss: 0.00290407496504
4856,	KL Loss: 1874.6441650390625	
It 27400: Total Loss: 6.031906604766846,		Rec Loss: 0.00294439122080
8029,	KL Loss: 1884.05078125	
It 27500: Total Loss: 6.041461944580078,		Rec Loss: 0.00325835496187
2101,	KL Loss: 1886.938720703125	
It 27600: Total Loss: 6.0324387550354,	Rec Loss: 0.0029912670142948627,	
	KL Loss: 1884.202392578125	
It 27700: Total Loss: 6.0444722175598145,		Rec Loss: 0.00308113498613
23833,	KL Loss: 1887.9346923828125	
It 27800: Total Loss: 6.091704368591309,		Rec Loss: 0.00359268812462
6875,	KL Loss: 1902.5350341796875	
It 27900: Total Loss: 6.084866046905518,		Rec Loss: 0.00303338211961
09056,	KL Loss: 1900.57275390625	
It 28000: Total Loss: 6.114672660827637,		Rec Loss: 0.00314617576077
58045,	KL Loss: 1909.85205078125	
It 28100: Total Loss: 6.101531028747559,		Rec Loss: 0.00298368977382
7791,	KL Loss: 1905.796142578125	
Run Epoch 30		
It 28200: Total Loss: 6.138578414916992,		Rec Loss: 0.00302832946181
2973,	KL Loss: 1917.3594970703125	
It 28300: Total Loss: 6.1455864906311035,		Rec Loss: 0.00306647038087
2488,	KL Loss: 1919.53759765625	
It 28400: Total Loss: 6.1561055183410645,		Rec Loss: 0.00339682307094
33556,	KL Loss: 1922.721435546875	
It 28500: Total Loss: 6.136283874511719,		Rec Loss: 0.00308250333182
5137,	KL Loss: 1916.62548828125	
It 28600: Total Loss: 6.202587127685547,		Rec Loss: 0.00316999573260
54573,	KL Loss: 1937.31787109375	
It 28700: Total Loss: 6.182847023010254,		Rec Loss: 0.00299033219926
05925,	KL Loss: 1931.205322265625	
It 28800: Total Loss: 6.226450443267822,		Rec Loss: 0.00296075944788
754,	KL Loss: 1944.840576171875	
It 28900: Total Loss: 6.240612983703613,		Rec Loss: 0.00293598882853
98483,	KL Loss: 1949.274169921875	
It 29000: Total Loss: 6.235678672790527,		Rec Loss: 0.00286587630398

```

5715,      KL Loss: 1947.7540283203125
Run Epoch 31
It 29100: Total Loss: 6.244664192199707,      Rec Loss: 0.00291323591955
00612,      KL Loss: 1950.54736328125
It 29200: Total Loss: 6.265792369842529,      Rec Loss: 0.00290708034299
31402,      KL Loss: 1957.151611328125
It 29300: Total Loss: 6.26191520690918,      Rec Loss: 0.00299979047849
77436,      KL Loss: 1955.9111328125
It 29400: Total Loss: 6.2388482093811035,      Rec Loss: 0.00285888137295
84217,      KL Loss: 1948.74658203125
It 29500: Total Loss: 6.270743370056152,      Rec Loss: 0.00304029113613
06906,      KL Loss: 1958.6572265625
It 29600: Total Loss: 6.274444103240967,      Rec Loss: 0.00265380553901
19553,      KL Loss: 1959.9345703125
It 29700: Total Loss: 6.288321018218994,      Rec Loss: 0.00346044008620
08333,      KL Loss: 1964.01904296875
It 29800: Total Loss: 6.30757474899292,      Rec Loss: 0.00320998579263
68713,      KL Loss: 1970.114013671875
It 29900: Total Loss: 6.318593502044678,      Rec Loss: 0.00256506889127
1949,      KL Loss: 1973.759033203125
It 30000: Total Loss: 6.315262794494629,      Rec Loss: 0.00257863500155
5085,      KL Loss: 1973.7188671875

```





Let's look at some reconstructions and decoded embedding samples!

```

In [ ]: # visualize VAE reconstructions and samples from the generative model
vis_reconstruction(vae_model)
vis_samples(vae_model)


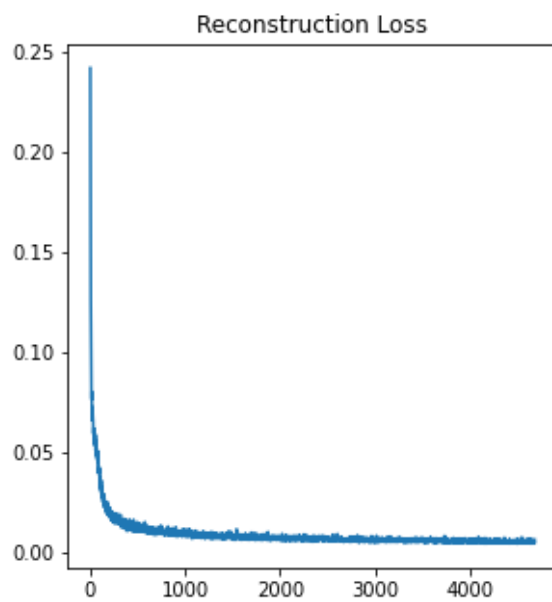
```






**Inline Question: What can you observe when setting  $\beta = 0$ ? Explain your observations! [3pt]** \ (please limit your answer to <150 words) \ **Answer:**

The model generates the original input almost precisely except some blurriness in the edges. I think it is almost similar to an autoencoder. For the generative part of the model where we sample from unit diagonal gaussian, the images generated are slightly more distinct and clear than the one generated by autoencoders. It seems as if the model is trying to generate some digits but most of them are deformed/skewed.



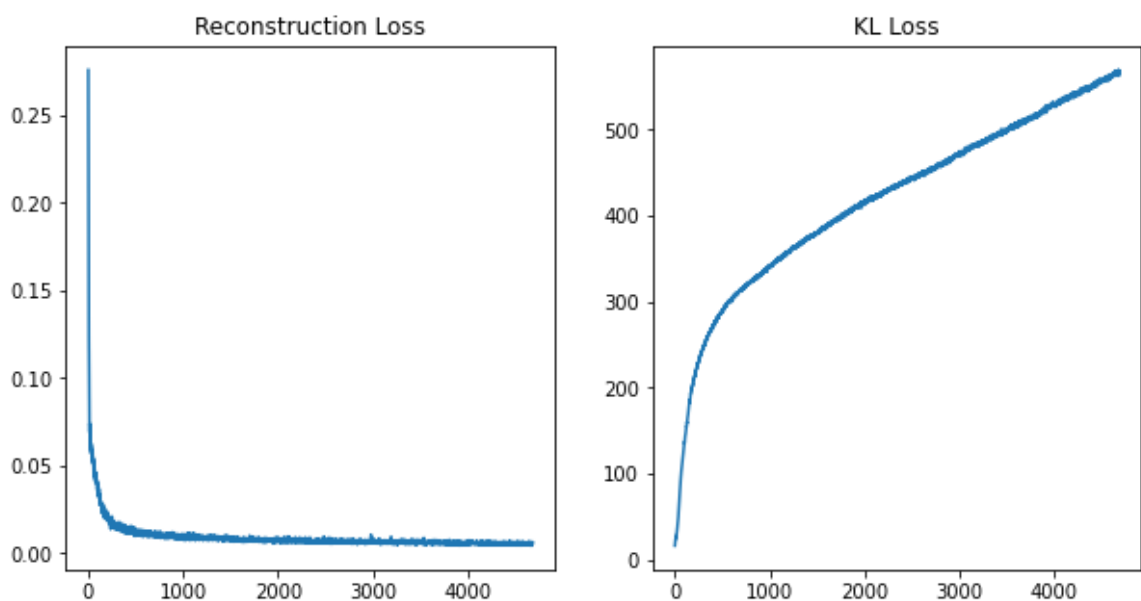




Let's repeat the same experiment for  $\beta = 10$ , a very high value for the coefficient. You can modify the  $\beta$  value in the cell above and rerun it (it is okay to overwrite the outputs of the previous experiment, but **make sure to copy the visualizations of training curves, reconstructions and samples for  $\beta = 0$  into your solution PDF** before deleting them).

**Inline Question: What can you observe when setting  $\beta = 10$ ?**

**Explain your observations! [3pt]** \ (please limit your answer to <200 words) \

**Answer:** It seems like at  $\beta=10$ , the digits in the reconstructed image are getting a little blurry. While the overall digits are almost the same, we can see additional pixels around the digits. Moreover, the samples of generation are also very blurry, resulting into almost undistinguishable digits. It seems that the digits that are generated are getting skewed in the space and not readable in many of the cases. It is like one digit is morphing into the next digit that is similar to it.





Now we can start tuning the beta value to achieve a good result. First describe what a "good result" would look like (focus what you would expect for reconstructions and sample quality).

**Inline Question: Characterize what properties you would expect for reconstructions (1pt) and samples (2pt) of a well-tuned VAE! [3pt]** \

(please limit your answer to <200 words) \ **Answer:** A well-tuned VAE should have sharper reconstructions and should have almost negligible blur around the digits. For the samples that are generated by a VAE, they should be more distinct and clear with respect to what digit they represent. Moreover, the skewness of the digits should not be too much that they are unreadable or unbelievable that a human would have written it.

## Tuning the $\beta$ -factor [5pt]

Now that you know what outcome we would like to obtain, try to tune  $\beta$  to achieve this result.

(logarithmic search in steps of 10x will be helpful, good results can be achieved after ~20 epochs of training). It is again okay to overwrite the results of the previous  $\beta=10$  experiment after copying them to the solution PDF.

## 4. Embedding Space Interpolation [3pt]

As mentioned in the introduction, AEs and VAEs cannot only be used to generate images, but also to learn low-dimensional representations of their inputs. In this final section we will investigate the representations we learned with both models by **interpolating in embedding space** between different images. We will encode two images into their low-dimensional embedding representations, then interpolate these embeddings and reconstruct the result.

```
In [ ]: START_LABEL = 5
        END_LABEL = 6
        nz=64

def get_image_with_label(target_label):
    """Returns a random image from the training set with the requested digit.
    for img_batch, label_batch in mnist_data_loader:
        for img, label in zip(img_batch, label_batch):
            if label == target_label:
                return img.to(device)

def interpolate_and_visualize(model, tag, start_img, end_img):
    """Encodes images and performs interpolation. Displays decodings."""
    model.eval() # put model in eval mode to avoid updating batchnorm

    # encode both images into embeddings (use posterior mean for interpolation)
    z_start = model.encoder(start_img[None])[..., :nz]
    z_end = model.encoder(end_img[None])[..., :nz]

    # compute interpolated latents
    N_INTER_STEPS = 5
    z_inter = [z_start + i/N_INTER_STEPS * (z_end - z_start) for i in range(N_INTER_STEPS)]

    # decode interpolated embeddings (as a single batch)
    img_inter = model.decoder(torch.cat(z_inter))

    # reshape result and display interpolation
    vis_imgs = torch.cat([start_img[None], img_inter, end_img[None]])
    fig = plt.figure(figsize = (10, 10))
    ax1 = plt.subplot(111)
    ax1.imshow(torchvision.utils.make_grid(vis_imgs, nrow=N_INTER_STEPS+2, padding=5,
                                           .data.cpu().numpy().transpose(1, 2, 0), cmap='gray'))
    plt.title(tag)
    plt.show()


# sample two training images with given labels
start_img = get_image_with_label(START_LABEL)
end_img = get_image_with_label(END_LABEL)

# visualize interpolations for AE and VAE models
interpolate_and_visualize(ae_model, "Auto-Encoder", start_img, end_img)
interpolate_and_visualize(vae_model, "Variational Auto-Encoder", start_img,
```



UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

cpuset\_checked))

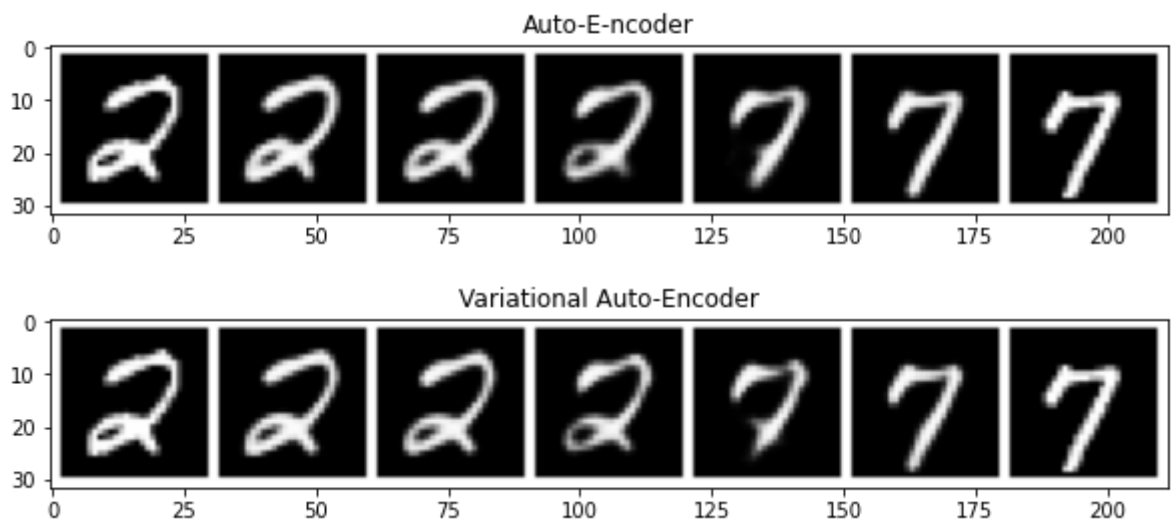


```
In [ ]: # sample two training images with given labels
start_img = get_image_with_label(2)
end_img = get_image_with_label(7)

# visualize interpolations for AE and VAE models
interpolate_and_visualize(ae_model, "Auto-E-ncoder", start_img, end_img)
interpolate_and_visualize(vae_model, "Variational Auto-Encoder", start_img,
```

/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:481:  
UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.


cpuset\_checked))



```
In [ ]: # sample two training images with given labels
start_img = get_image_with_label(3)
end_img = get_image_with_label(8)

# visualize interpolations for AE and VAE models
interpolate_and_visualize(ae_model, "Auto-E-ncoder", start_img, end_img)
interpolate_and_visualize(vae_model, "Variational Auto-Encoder", start_img,
```

/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:481: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.  
cpuset\_checked))



Repeat the experiment for different start / end labels and different samples. Describe your observations.

**Inline Question: Repeat the interpolation experiment with different start / end labels and multiple samples. Describe your observations! Focus on: \**

1. How do AE and VAE embedding space interpolations differ? \
2. How do you expect these differences to affect the usefulness of the learned representation for downstream learning? \ (please limit your answer to <300 words)

**Answer:** The interpolations of AE have more randomness compared to the embedding space interpolations from VAE. Comparing both the visualizations we can see that the latent space of VAE has more inclination to the MNIST like characters while AE is very blurry. Using AE and VAE in downstream learning will provide drastically different results because a random value in case of VAE generates more meaningful output than AE in somecases where is feels like AE is trying to morph one character into another or the latent space is overlapping.

Submission PDF

As in assignment 1, please prepare a separate submission PDF for each problem. **Do not simply render your notebook as a pdf.** For this problem, please include the following plots & answers in a PDF called `problem_1_solution.pdf` :

1. Auto-encoder samples and AE sampling inline question answer.
2. VAE training curves, reconstructions and samples for:
  - $\beta = 0$
  - $\beta = 10$
  - your tuned  $\beta$  (also listing the tuned value for  $\beta$ )
3. Answers to all inline questions in VAE section (ie 4 inline questions).
4. Three representative interpolation comparisons that show AE and VAE embedding interpolation between the same images.
5. Answer to interpolation inline question.

Note that you still need to submit the jupyter notebook with all generated solutions. We will randomly pick submissions and check that the plots in the PDF and in the notebook are equivalent (except for those  $\beta=0$  and  $\beta=10$  plots that we allowed to overwrite).